

Senior Data Engineer | Cloud & Data Platform Engineering (AWS, Python, SQL)Okotoks, AB, Canada | apiskun@gmail.com | +1 (431) 588-3044<https://www.linkedin.com/in/alexanderpiskun/> | <https://github.com/AlexanderPiskun>**Professional Summary**

Senior Data Engineer with 10+ years of experience architecting mission-critical ground-up ETL/ELT solutions for global enterprises. Expert in Python and advanced SQL, specializing in processing massive datasets within AWS and Snowflake ecosystems to deliver high-performance, cost-efficient data products. Proven leader in end-to-end pipeline ownership, from ingestion and transformation to automated monitoring and data reliability engineering. A strategic partner to Product Managers and Data Analysts, translating complex business requirements into actionable, high-quality data products within agile, fully remote environments.

Core Competencies & Technical Stack**Data Engineering & Pipelines**

- ETL/ELT pipeline design and development (batch, near real-time, and event-driven)
- Large-scale data processing using Python, SQL, and dbt
- Data ingestion, transformation, and orchestration (Airflow concepts)

Cloud & Infrastructure

- AWS: S3, Lambda, EC2, RDS, Glue, Step Functions, SQS, SNS
- GCP: GCP Cloud Storage
- Infrastructure as Code: Terraform, AWS CloudFormation

Data Modeling & Warehousing

- Relational data modeling, normalization, performance optimization
- Data warehouse design (Snowflake, Redshift)

Software Engineering for Data

- Unit Testing (PyTest), CI/CD (GitHub Actions), Version Control (Git), Documentation, and Data Reliability Engineering

Security & Governance

- Data Privacy, RBAC, Data Masking (Dynamic & Tag-based), Row-Level Security, and SOC2/PII compliance; experienced operating in regulated data environments with audit evidence requirements (insurance, financial services)

Data Quality

- Automated validation frameworks, Snowflake Data Metric Functions (DMFs), and data observability & monitoring

Work Experience**Database Administrator & Data Engineer****Feb 2019 - Oct 2025**

Wawanesa Mutual Insurance Company, Canada

Impact Summary:

Delivered highly reliable, cost-efficient cloud data platforms supporting mission-critical financial workloads, improving performance, scalability, and operational efficiency across AWS and Snowflake environments.

- Led technical discovery and solution design for data integration initiatives, translating business requirements into scalable, production-ready data pipelines
- Designed and implemented scalable ETL pipelines using Python and AWS services, enabling automated, reliable data ingestion and transformation across high-volume transactional systems processing millions of records daily
- Optimized data pipelines and query performance for large-scale workloads, improving execution efficiency and reducing processing time for analytics use cases by ~70%
- Owned data pipeline reliability and observability, implementing unit testing, monitoring, real-time alerting, CI/CD, and automated data quality validation to ensure trusted data for downstream consumers

- Improved data pipeline efficiency and reduced infrastructure costs by optimizing workloads and implementing scalable ingestion patterns across AWS and Snowflake environments
- Architected and maintained enterprise-scale on-prem, Oracle Exadata Cloud@Customer (ExaCC), and AWS RDS database environments supporting financial and insurance data platforms, ensuring excellence of operation and 99.99% uptime for core data
- Took ownership of data platform architecture, including ingestion, transformation, storage, and delivery layers, supporting analytics and business decision-making
- Implemented rigorous RBAC, encryption (Data-at-Rest & Data-in-Transit), and access controls to meet strict compliance requirements (SOC2/PII), mirroring fintech and insurance standards
- Implemented automated data quality checks (freshness, row count, consistency) using Snowflake Data Metric Functions, aligning with dbt-style testing best practices
- Collaborated with product managers, analytics, data science, and cross-functional teams to deliver scalable data solutions and support downstream analytics and reporting requirements

Senior Database Administrator & Data Engineer**Jan 2006 - Jun 2018**

Yael Group, Israel

Impact Summary:

Transformed legacy database environments into scalable, high-performance data platforms through automation, optimization, and robust ETL design, significantly improving system reliability and reporting efficiency.

- Architected scalable data pipelines using MS SQL SSIS, Oracle XML DB, and custom ETL modules to unify disparate data sources for downstream analytics
- Optimized complex OLTP and Data Warehouse environments, reducing long-running report execution times from hours to seconds through advanced SQL tuning and schema refactoring
- Developed a proactive Python/Shell-based monitoring framework that significantly decreased unplanned downtime and transitioned the team from reactive to predictive maintenance
- Designed and implemented early-stage ETL pipelines supporting data warehousing and reporting initiatives
- Designed and deployed a custom Windows-based automation service to ingest and process external data attachments, eliminating manual entry and establishing a standardized data interface for the organization
- Built a comprehensive library of automation scripts (PowerShell, VBS, Perl, Java) to streamline administrative overhead and standardize cross-platform database operations
- Orchestrated full-lifecycle administration, patching, and multi-version upgrades for Oracle (9i–12c) and MS SQL Server (2012/2014) ecosystems, ensuring 100% alignment with project milestones and enterprise security standards

Key Projects**Selected Project: Cloud Data Platform (Snowflake + Python + dbt)**

- Designed and implemented a scalable cloud data platform supporting ingestion, transformation, and analytics workflows
- Built SQL and Python-based ETL pipelines to process and transform large datasets into analytics-ready formats
- Designed and deployed cross-cloud data connectors to ingest raw objects from AWS S3 and GCP Cloud Storage, leveraging Python and Snowflake-native integration features to reduce data silos
- Implemented data quality validation, monitoring, and automated testing to ensure reliability
- Delivered optimized data models and data marts for downstream analytics and reporting
- Applied CI/CD and version control practices to manage data pipeline deployments
- Designed a clean RBAC model separating engineering, analytics, and consumer personas to enforce least-privilege access
- Built incremental processing pipelines, reducing processing costs by approximately 30% while improving pipeline reliability

Education & Certifications**Bachelor of Science (BS), Mathematics & Computer Science | Ben-Gurion University of the Negev, Israel**

- **Snowflake SnowPro Core Certification (Snowflake)**
Credential ID/Verification: <https://achieve.snowflake.com/ed4a74b7-639b-4981-ac39-7ce049459d24>
- **AWS Certified Solutions Architect – Associate (In Progress)**